

Max Time: 2.5 hours

Date: 15-02-2023

**Instructions:**

- Please provide your own solutions and DO NOT COPY the code from your colleagues or the web.
- You can discuss your problems only with the teachers.
- All tasks carry equal points.

**Task # 01 - Linear Regression**

**[8 x 5 = 40]**

**Problem Statement**

Write a Python function `get_diabetes_predictions(feature_idx)` that takes an integer `feature_idx` as input, selects the input feature at index `feature_idx` from the diabetes dataset, trains a linear regression model using the selected feature, makes predictions on the testing set, and returns the model, mean squared error and coefficient of determination of the model on the testing set.

**Signature**

```
def get_diabetes_predictions(feature_idx:int) -> Tuple[LinearRegression, float, float]:  
    pass
```

**Input**

- `feature_idx` ( $1 \leq \text{feature\_idx} \leq 10$ ): An integer representing the index of the input feature to use for training the model.

**Output**

- A tuple containing the following elements:
  - `LinearRegression`: The trained linear regression model.
  - `float`: The mean squared error of the model on the testing set.
  - `float`: The coefficient of determination (R-squared score) of the model on the testing set.

## Instructions

Your program should perform the following steps:

1. Load the diabetes dataset using Scikit-learn's `datasets.load_diabetes` function.
2. Split the dataset into training and testing sets using scikit-learn's `train_test_split` function. Use 80% of the data for training and 20% for testing.
3. Select the input feature at index `feature_idx` from the dataset.
4. Create a linear regression model using scikit-learn's `LinearRegression` class.
5. Fit the model to the training data using the `fit` method.
6. Make predictions on the testing data using the `predict` method.
7. Calculate the mean squared error (MSE) and R-squared score of the model on the testing data using scikit-learn's `mean_squared_error` and `r2_score` functions.
8. Return the trained model, MSE and R-squared score as a tuple.

## Example Usage

```
model, mse, r2 = get_diabetes_predictions(2)
print("Mean squared error: {:.2f}".format(mse))
print("R-squared score: {:.2f}".format(r2))
```

## Task # 02 - MLP Classification

[7x5 = 35]

### Problem Statement

Write a Python function `get_diabetes_mlp_predictions(hidden_layers)` that takes a tuple `hidden_layers` as input, representing the number of hidden layers and number of neurons in each hidden layer of an `MLPClassifier`. The function should train the `MLPClassifier` on the diabetes dataset, make predictions on the testing set, and return the model, mean squared error and coefficient of determination of the model on the testing set.

### Signature

```
def get_diabetes_mlp_predictions(hidden_layers: Tuple[int]) -> Tuple[MLPClassifier, float, float]:  
    pass
```

### Input

- `hidden_layers` ( $1 \leq \text{len}(\text{hidden\_layers}) \leq 5$ ): A tuple representing the number of hidden layers and number of neurons in each hidden layer of the `MLPClassifier`. The tuple should contain between 1 and 5 integers, where each integer represents the number of neurons in that hidden layer. For example, (10, 5) would represent an `MLPClassifier` with 2 hidden layers, the first with 10 neurons and the second with 5 neurons.

### Output

- A tuple containing the following elements:
  - `MLPClassifier`: The trained `MLPClassifier` model.
  - float: The mean squared error of the model on the testing set.
  - float: The coefficient of determination (R-squared score) of the model on the testing set.

### Instructions

Your program should perform the following steps:

1. Load the diabetes dataset using Scikit-learn's `datasets.load_diabetes` function.
2. Split the dataset into training and testing sets using scikit-learn's `train_test_split` function. Use 80% of the data for training and 20% for testing.
3. Create an `MLPClassifier` using scikit-learn's `MLPClassifier` class. Use the

hidden\_layers tuple to specify the number of hidden layers and number of neurons in each hidden layer.

4. Fit the model to the training data using the fit method.
5. Make predictions on the testing data using the predict method.
6. Calculate the mean squared error (MSE) and R-squared score of the model on the testing data using scikit-learn's mean\_squared\_error and r2\_score functions.
7. Return the trained model, MSE and R-squared score as a tuple.

### **Example Usage**

```
model, mse, r2 = get_diabetes_mlp_predictions((10, 5))
print("Mean squared error: {:.2f}".format(mse))
print("R-squared score: {:.2f}".format(r2))
```